

Binary Prediction with a Rainfall Dataset

April 5, 2025

1 Binary Prediction with a Rainfall Dataset

1.0.1 Importing and loading necessary libraries and packages

```
[1]: import math
import warnings
from itertools import combinations
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_theme(style='white', palette='Set2')
pal = sns.color_palette('Set2')
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import StackingClassifier
from sklearn.pipeline import Pipeline
from sklearn.model_selection import
    ↪(train_test_split,StratifiedKFold,cross_val_score)
from sklearn.metrics import (roc_curve,roc_auc_score,classification_report,auc)
from sklearn.preprocessing import MinMaxScaler
from sklearn.utils.class_weight import compute_class_weight
import xgboost as xgb
from xgboost import XGBClassifier
xgb.set_config(verbosity=0)
import lightgbm
from lightgbm import LGBMClassifier, plot_importance
import optuna
from optuna.samplers import TPESampler
import logging
optuna.logging.set_verbosity(logging.WARNING)
import warnings
warnings.filterwarnings("ignore", category=RuntimeWarning)
warnings.filterwarnings("ignore")
```

1.0.2 Loading Data, Data Preprocessing, and Visualizing Data

```
[2]: file_path = '/Users/nav/Documents/Dataset/Rainfall Dataset/train.csv'  
file_path1 = '/Users/nav/Documents/Dataset/Rainfall Dataset/test.csv'
```

```
[3]: train_df = pd.read_csv(file_path)  
test_df = pd.read_csv(file_path1)
```

```
[4]: train_df.head()
```

```
[4]:    id  day  pressure  maxtemp  temperature  mintemp  dewpoint  humidity  \  
0   0    1     1017.4    21.2       20.6      19.9      19.4      87.0  
1   1    2     1019.5    16.2       16.9      15.8      15.4      95.0  
2   2    3     1024.1    19.4       16.1      14.6      9.3       75.0  
3   3    4     1013.4    18.1       17.8      16.9      16.8      95.0  
4   4    5     1021.8    21.3       18.4      15.2      9.6       52.0  
  
    cloud  sunshine  winddirection  windspeed  rainfall  
0   88.0      1.1        60.0      17.2        1  
1   91.0      0.0        50.0      21.9        1  
2   47.0      8.3        70.0      18.1        1  
3   95.0      0.0        60.0      35.6        1  
4   45.0      3.6        40.0      24.8        0
```

```
[5]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 2190 entries, 0 to 2189  
Data columns (total 13 columns):  
 #  Column          Non-Null Count  Dtype     
---  --    
 0   id              2190 non-null   int64    
 1   day             2190 non-null   int64    
 2   pressure        2190 non-null   float64  
 3   maxtemp         2190 non-null   float64  
 4   temperature     2190 non-null   float64  
 5   mintemp         2190 non-null   float64  
 6   dewpoint        2190 non-null   float64  
 7   humidity         2190 non-null   float64  
 8   cloud            2190 non-null   float64  
 9   sunshine         2190 non-null   float64  
 10  winddirection   2190 non-null   float64  
 11  windspeed        2190 non-null   float64  
 12  rainfall         2190 non-null   int64    
dtypes: float64(10), int64(3)  
memory usage: 222.6 KB
```

```
[6]: train_df.describe()
```

```
[6]:
```

	id	day	pressure	maxtemp	temparature	\
count	2190.000000	2190.000000	2190.000000	2190.000000	2190.000000	
mean	1094.500000	179.948402	1013.602146	26.365799	23.953059	
std	632.342866	105.203592	5.655366	5.654330	5.222410	
min	0.000000	1.000000	999.000000	10.400000	7.400000	
25%	547.250000	89.000000	1008.600000	21.300000	19.300000	
50%	1094.500000	178.500000	1013.000000	27.800000	25.500000	
75%	1641.750000	270.000000	1017.775000	31.200000	28.400000	
max	2189.000000	365.000000	1034.600000	36.000000	31.500000	
	mintemp	dewpoint	humidity	cloud	sunshine	\
count	2190.000000	2190.000000	2190.000000	2190.000000	2190.000000	
mean	22.170091	20.454566	82.036530	75.721918	3.744429	
std	5.059120	5.288406	7.800654	18.026498	3.626327	
min	4.000000	-0.300000	39.000000	2.000000	0.000000	
25%	17.700000	16.800000	77.000000	69.000000	0.400000	
50%	23.850000	22.150000	82.000000	83.000000	2.400000	
75%	26.400000	25.000000	88.000000	88.000000	6.800000	
max	29.800000	26.700000	98.000000	100.000000	12.100000	
	winddirection	windspeed	rainfall			
count	2190.000000	2190.000000	2190.000000			
mean	104.863151	21.804703	0.753425			
std	80.002416	9.898659	0.431116			
min	10.000000	4.400000	0.000000			
25%	40.000000	14.125000	1.000000			
50%	70.000000	20.500000	1.000000			
75%	200.000000	27.900000	1.000000			
max	300.000000	59.500000	1.000000			

```
[7]: train_df.isna().sum()
```

```
[7]:
```

	id	0
day	0	
pressure	0	
maxtemp	0	
temparature	0	
mintemp	0	
dewpoint	0	
humidity	0	
cloud	0	
sunshine	0	
winddirection	0	
windspeed	0	
rainfall	0	
dtype:	int64	

```
[8]: test_df.head()
```

```
[8]:      id  day  pressure  maxtemp  temperature  mintemp  dewpoint  humidity \
0   2190     1    1019.5     17.5       15.8      12.7      14.9      96.0
1   2191     2    1016.5     17.5       16.5      15.8      15.1      97.0
2   2192     3    1023.9     11.2       10.4      9.4       8.9      86.0
3   2193     4    1022.9     20.6       17.3      15.2      9.5      75.0
4   2194     5    1022.2     16.1       13.8      6.4       4.3      68.0

      cloud  sunshine  winddirection  windspeed
0    99.0      0.0        50.0      24.3
1    99.0      0.0        50.0      35.3
2    96.0      0.0        40.0      16.9
3    45.0      7.1        20.0      50.6
4    49.0      9.2        20.0      19.4
```

```
[9]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 12 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                730 non-null    int64  
 1   day               730 non-null    int64  
 2   pressure          730 non-null    float64 
 3   maxtemp           730 non-null    float64 
 4   temperature       730 non-null    float64 
 5   mintemp           730 non-null    float64 
 6   dewpoint          730 non-null    float64 
 7   humidity          730 non-null    float64 
 8   cloud              730 non-null    float64 
 9   sunshine          730 non-null    float64 
 10  winddirection     729 non-null    float64 
 11  windspeed         730 non-null    float64 
dtypes: float64(10), int64(2)
memory usage: 68.6 KB
```

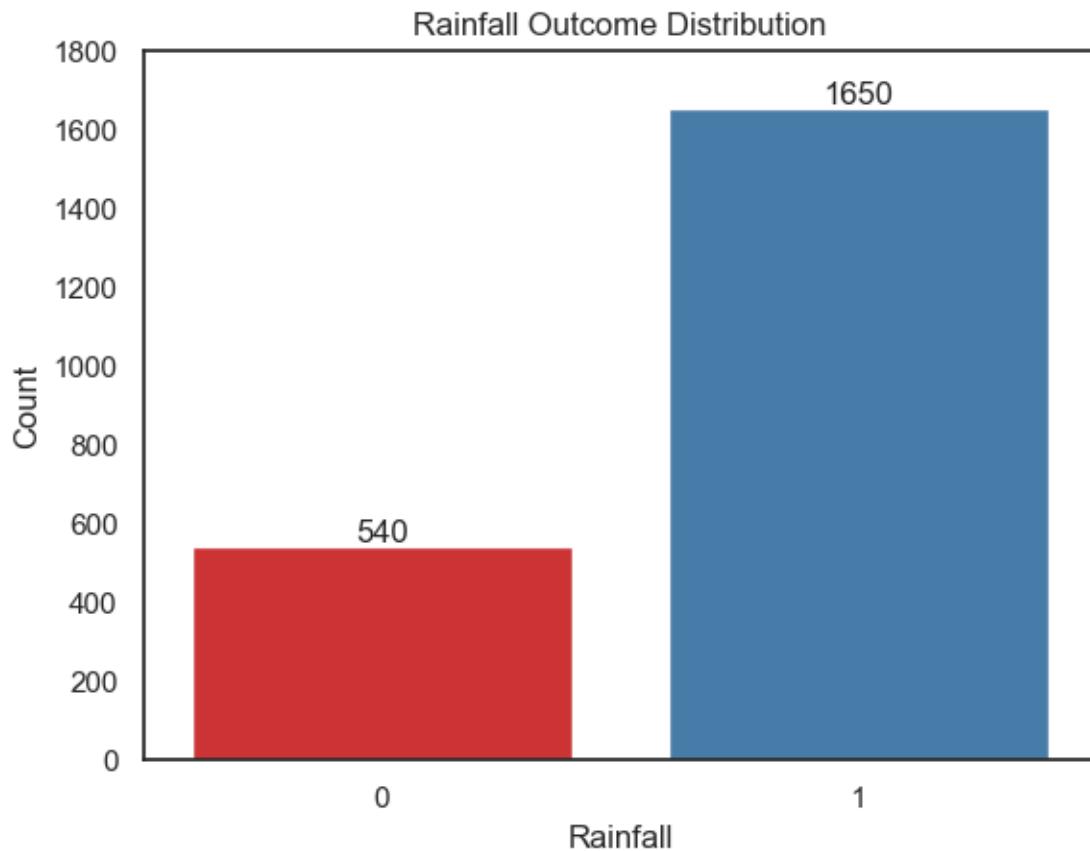
```
[10]: test_df.isna().sum()
```

```
[10]: id                0
      day               0
      pressure          0
      maxtemp           0
      temperature       0
      mintemp           0
      dewpoint          0
      humidity          0
```

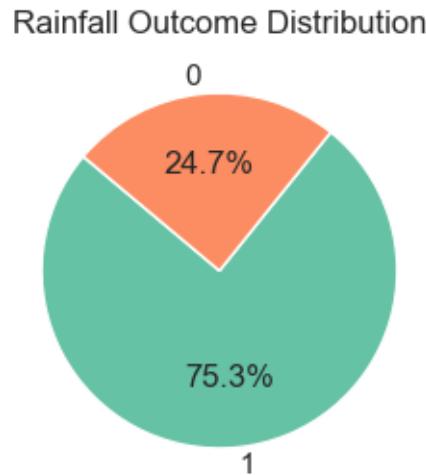
```
cloud          0  
sunshine       0  
winddirection  1  
windspeed      0  
dtype: int64
```

```
[11]: # Fill the missing value in 'winddirection' with the median (since it's unidirectional)  
test_df['winddirection'] = test_df['winddirection'].  
   fillna(test_df['winddirection'].median())
```

```
[12]: ax = sns.countplot(x='rainfall', data=train_df, palette='Set1')  
for container in ax.containers:  
    ax.bar_label(container)  
  
ax.set_ylabel('Count')  
ax.set_xlabel('Rainfall')  
ax.set_title('Rainfall Outcome Distribution')  
ax.set_ylim(0, 1800)  
plt.show()
```

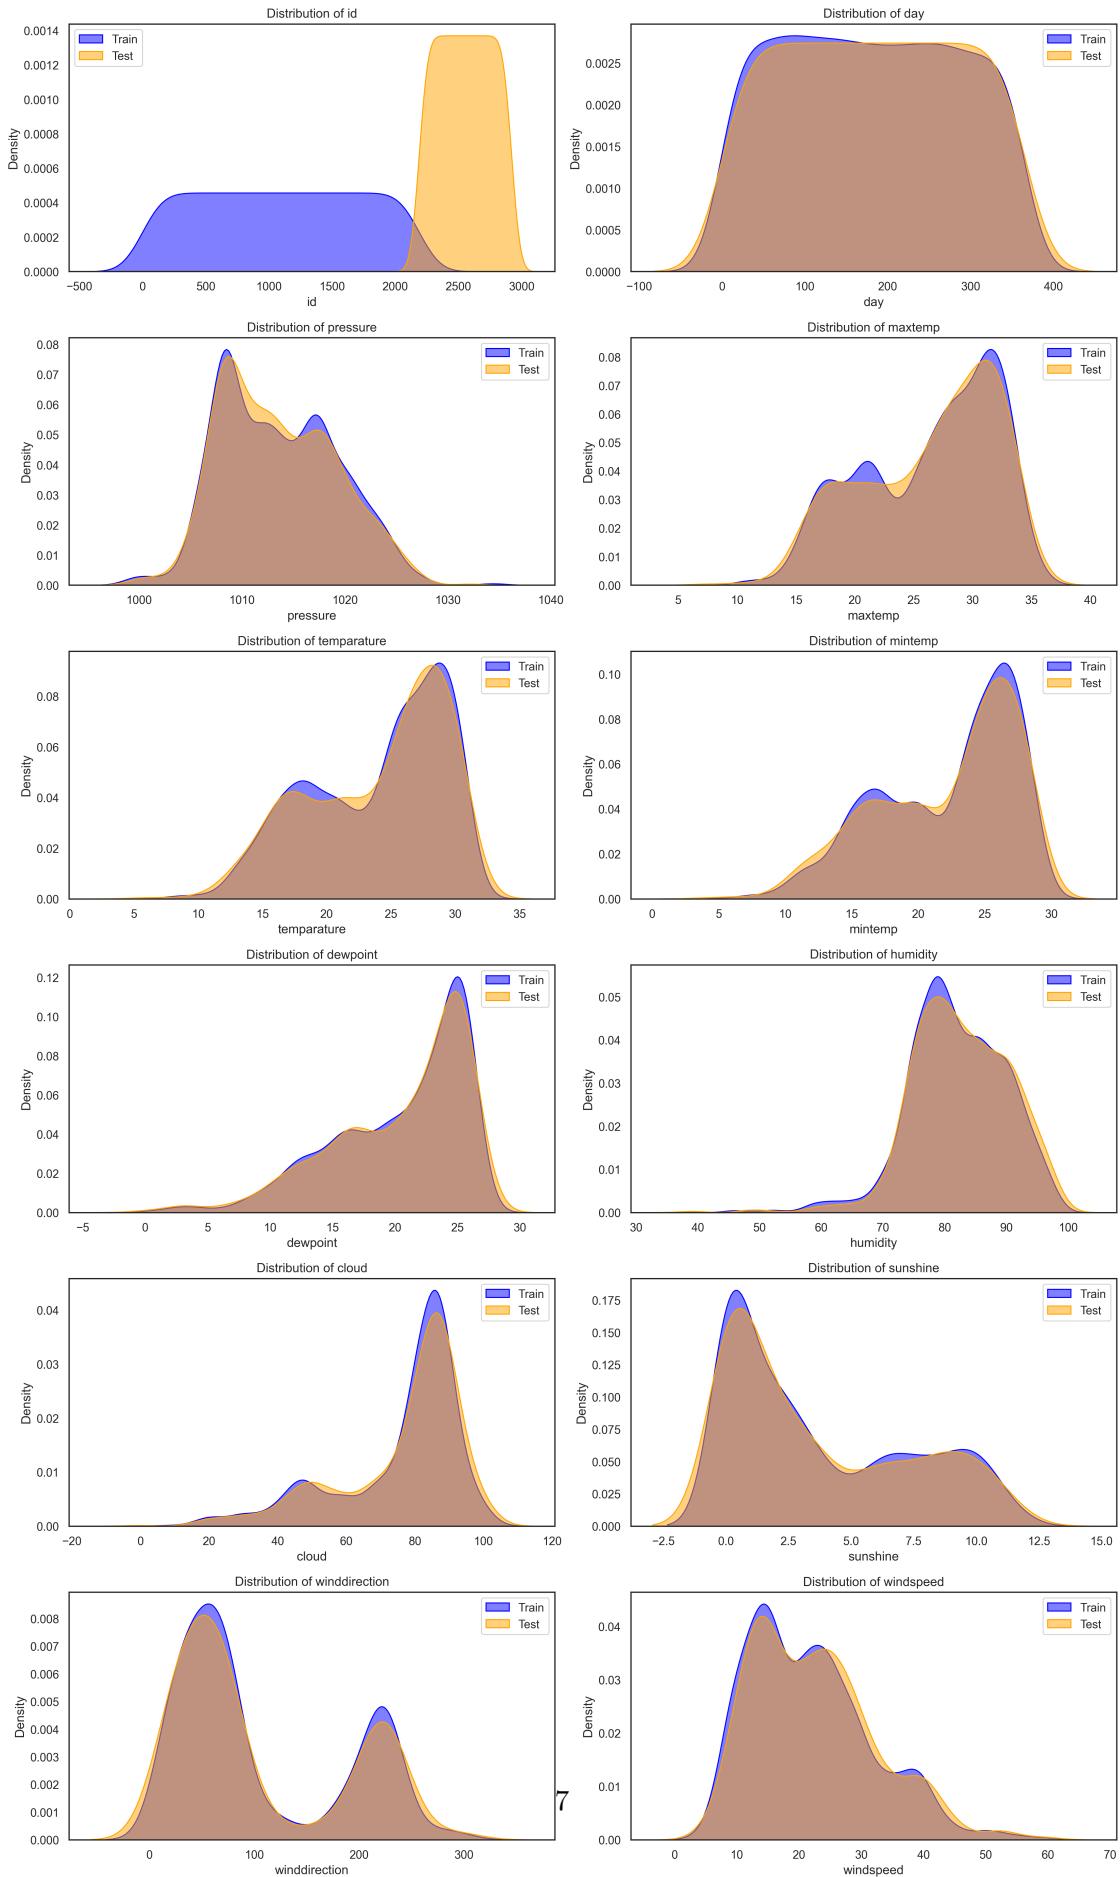


```
[13]: data = train_df['rainfall'].value_counts()
plt.figure(figsize=(3, 3))
plt.pie(data, labels=data.index, autopct='%.1f%%', startangle=140)
plt.title("Rainfall Outcome Distribution")
plt.show()
```



```
[14]: # Define features
features = ['id', 'day', 'pressure', 'maxtemp', 'temparature', 'mintemp',
            'dewpoint', 'humidity', 'cloud', 'sunshine', 'winddirection', ↴
            'windspeed']
fig, ax = plt.subplots(6, 2, figsize=(15, 25), dpi=300)
ax = ax.flatten()

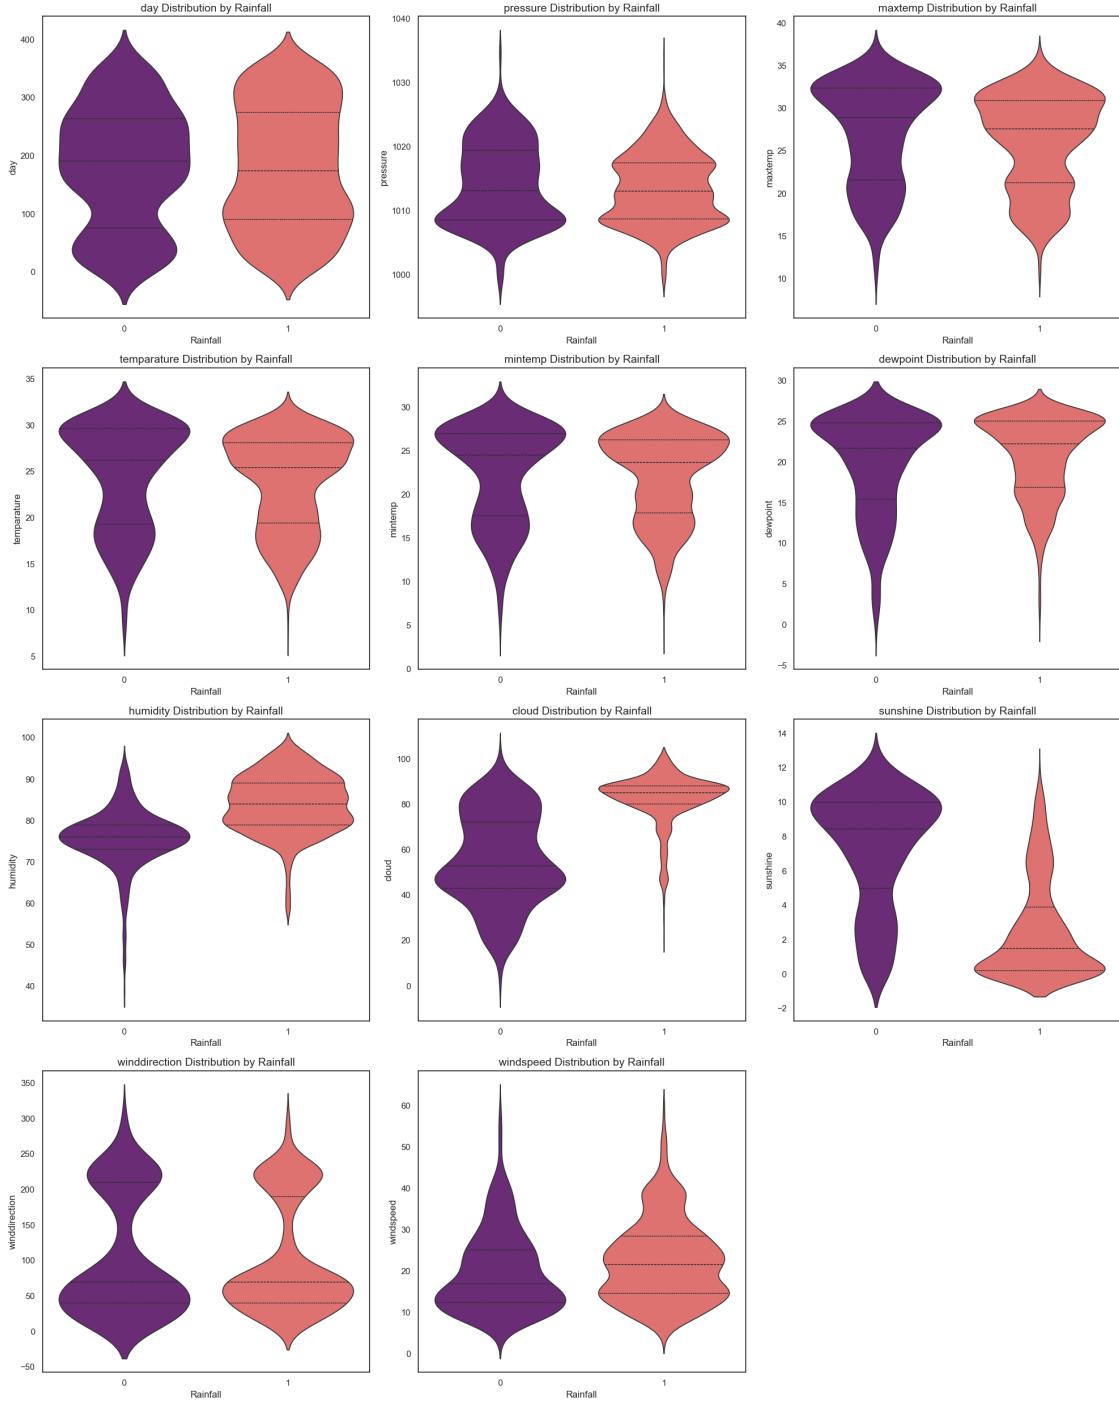
# Plot KDE for both train_df and test_df
for i, feature in enumerate(features):
    sns.kdeplot(train_df[feature], ax=ax[i], fill=True, label='Train', ↴
                color='blue', alpha=0.5)
    sns.kdeplot(test_df[feature], ax=ax[i], fill=True, label='Test', ↴
                color='orange', alpha=0.5)
    ax[i].set_title(f'Distribution of {feature}')
    ax[i].set_xlabel(feature)
    ax[i].set_ylabel('Density')
    ax[i].legend()
plt.tight_layout()
plt.show()
```



The histograms above illustrate the distribution of each numerical feature for both the training and testing datasets. By comparing the two, we can observe that the train and test sets exhibit similar distribution patterns, indicating that the data split is consistent and representative.

```
[15]: features = [col for col in train_df.columns if col not in ['id', 'rainfall']]
plt.figure(figsize=(20, 25))
for i, col in enumerate(features, 1):
    plt.subplot(4, 3, i)
    sns.violinplot(data=train_df, x='rainfall', y=col, palette='magma', inner="quartile", linewidth=1.2)
    plt.title(f'{col} Distribution by Rainfall', fontsize=14)
    plt.xlabel('Rainfall', fontsize=12)
    plt.ylabel(col, fontsize=12)

plt.tight_layout()
plt.show()
```



The violin plots above visualize the relationship between the target variable ‘rainfall’ and features such as ‘humidity’, ‘cloud’, and ‘sunshine’, providing deeper insight into their distributions.

- **Rainfall vs Humidity:** Class 0 (no rainfall) tends to cluster around 75% humidity, while class 1 (rainfall) clusters closer to 85%. This suggests that rainfall is more likely when humidity levels are higher.

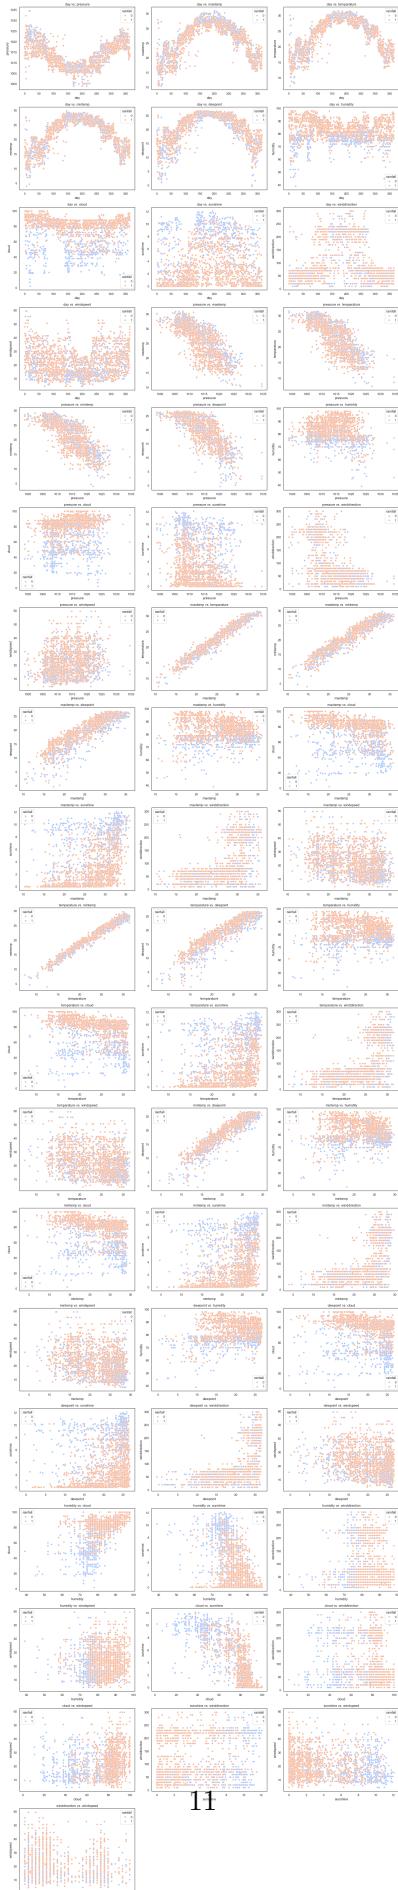
- **Rainfall vs Cloud Coverage:** Class 0 is centered around 55% cloud cover, whereas class 1 clusters around 85%, indicating that increased cloud coverage may be associated with a higher likelihood of rainfall.
- **Rainfall vs Sunshine:** Class 0 tends to have around 8 hours of sunshine, while class 1 is closer to just 1 hour. This implies that rainfall typically occurs on days with significantly less sunshine.

```
[16]: # Define features and create feature pairs
features = ['day', 'pressure', 'maxtemp', 'temparature', 'mintemp',
            'dewpoint', 'humidity', 'cloud', 'sunshine', 'windirection', ↴
            'windspeed']
pairs = list(combinations(features, 2))
n_cols = 3
n_rows = math.ceil(len(pairs) / n_cols)
fig, axes = plt.subplots(n_rows, n_cols, figsize=(20, 5 * n_rows))
axes = axes.flatten()

for i, (x, y) in enumerate(pairs):
    sns.scatterplot(data=train_df, x=x, y=y, hue='rainfall', ↴
                    palette='coolwarm', ax=axes[i])
    axes[i].set_title(f'{x} vs. {y}')

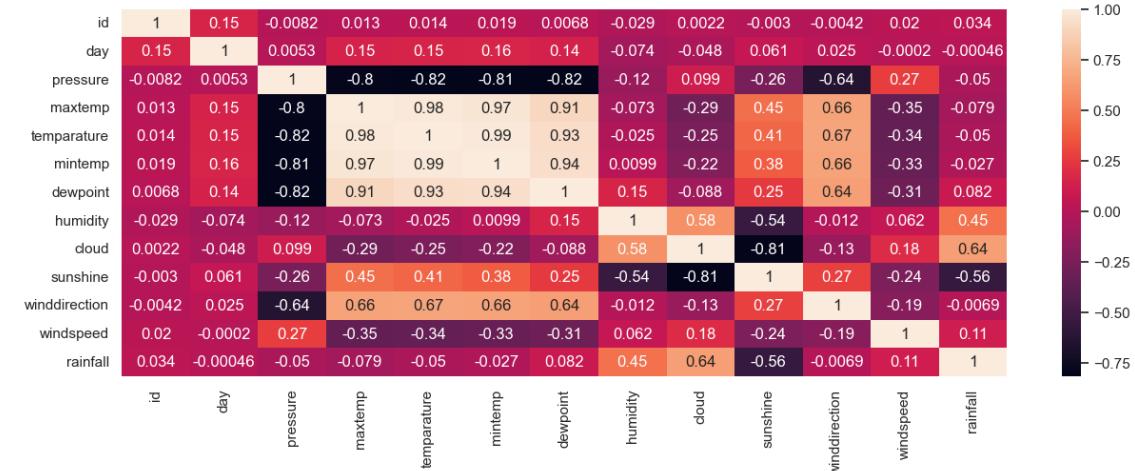
for j in range(i + 1, len(axes)):
    axes[j].axis('off')

plt.tight_layout()
plt.show()
```



```
[17]: plt.figure(figsize = (15,5))
sns.heatmap(train_df.corr(), annot = True)
plt.show
```

[17]: <function matplotlib.pyplot.show(close=None, block=None)>



```
[18]: target_column = train_df.columns[-1]
target_corr = train_df.corr()[target_column].sort_values(ascending=False)
print(target_corr.to_frame())
```

	rainfall
rainfall	1.000000
cloud	0.641191
humidity	0.454213
windspeed	0.111625
dewpoint	0.081965
id	0.033674
day	-0.000462
winddirection	-0.006939
mintemp	-0.026841
temparature	-0.049660
pressure	-0.049886
maxtemp	-0.079304
sunshine	-0.555287

- Positive Correlations:

- Some features show strong positive relationships (correlation values near +1), meaning they tend to increase together.
- For example, **Cloud (0.641)** and **Humidity (0.454)** show moderately strong positive

correlations with the target variable **Rainfall**, indicating that higher cloud coverage and humidity are generally associated with more rainfall.

- **Negative Correlations:**
 - Strong negative values (close to -1) suggest an inverse relationship, where one variable increases as the other decreases.
- **Low or No Correlation:**
 - Features like **Windspeed (0.112)** and **Dewpoint (0.082)** show very weak positive correlations with Rainfall, indicating minimal predictive value.
- **Multicollinearity Risk:**
 - High correlations between independent variables could indicate multicollinearity, which may affect the stability and interpretability of predictive models.
- **Cloud coverage and humidity** are the most influential factors in predicting rainfall.
- **Wind speed** and **dewpoint** show weak relationships and may contribute less to model performance.
- When building predictive models, consider prioritizing features with stronger correlations to the target.

1.0.3 Feature Engineering

```
[19]: # create the feature engineering function
def create_features(df):
    df_new = df.copy()

    df_new['temp_range'] = df_new['maxtemp'] - df_new['mintemp']
    df_new['temp_ratio'] = df_new['temparature'] / (df_new['maxtemp'] + 1e-5)
    df_new['temp_from_dewpoint'] = df_new['temparature'] - df_new['dewpoint']
    df_new['max_min_temp_ratio'] = df_new['maxtemp'] / (df_new['mintemp'] +
    ↪1e-5)

    df_new['humid_temp_interaction'] = df_new['humidity'] * ↪
    ↪df_new['temparature']
    df_new['humid_pressure_interaction'] = df_new['humidity'] * ↪
    ↪df_new['pressure']
    df_new['humid_dewpoint_interaction'] = df_new['humidity'] * ↪
    ↪df_new['dewpoint']

    df_new['cloud_sunshine_ratio'] = df_new['cloud'] / (df_new['sunshine'] + 1)
    df_new['cloud_coverage_rate'] = df_new['cloud'] / 100

    df_new['pressure_temp_ratio'] = df_new['pressure'] / (df_new['temparature'] +
    ↪1e-5)
    df_new['pressure_humidity_ratio'] = df_new['pressure'] / ↪
    ↪(df_new['humidity'] + 1e-5)

    df_new['weather_severity'] = (df_new['cloud'] * df_new['humidity']) / ↪
    ↪(df_new['pressure'] * (df_new['sunshine'] + 1))
```

```

df_new['temp_humidity_index'] = (df_new['temparature'] * df_new['humidity']) / 100
df_new['pressure_temp_humidity'] = (df_new['pressure'] * df_new['temparature']) / (df_new['humidity'] + 1e-5)

df_new['day_sin'] = np.sin(2 * np.pi * df_new['day'] / 365)
df_new['day_cos'] = np.cos(2 * np.pi * df_new['day'] / 365)

# Wind feature
df_new['wind_power'] = df_new['windspeed'] * np.abs(np.cos(np.radians(df_new['winddirection'])))
```

return df_new

Apply feature engineering

train_fe = create_features(train_df)

Prepare data for feature selection

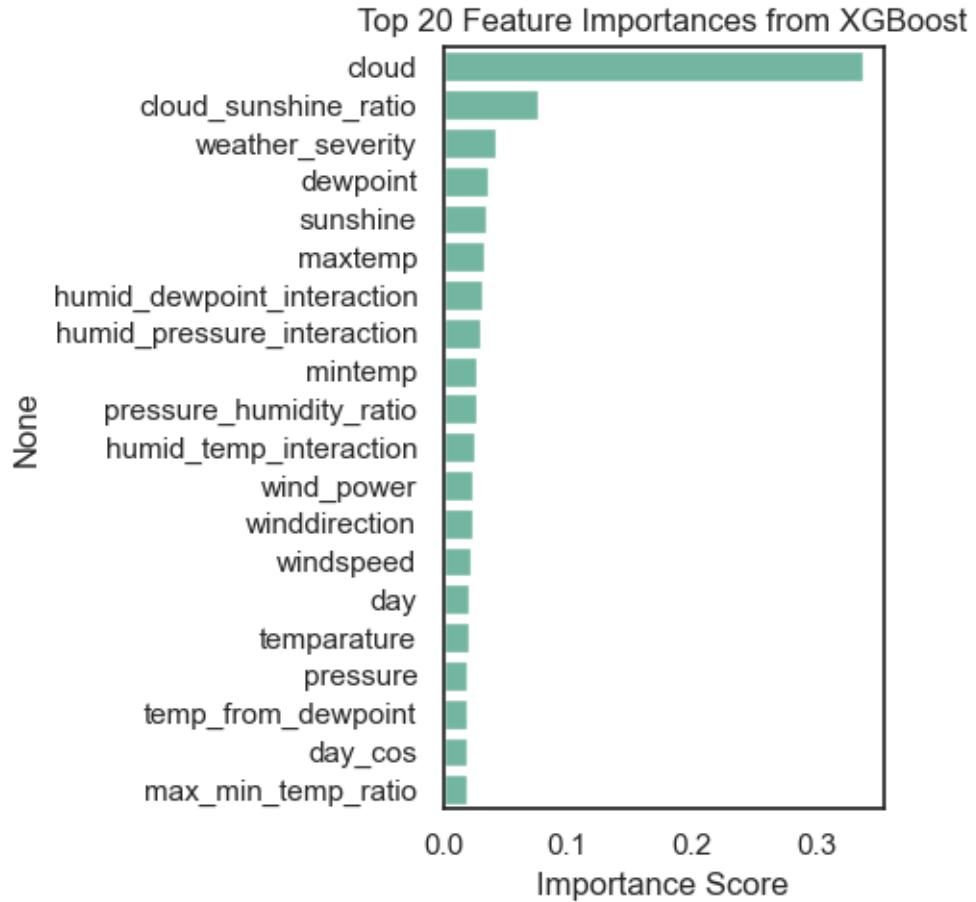
X = train_fe.drop(columns=['id', 'rainfall'])
y = train_fe['rainfall']
X_train, X_val, y_train, y_val = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)

Train XGBoost model

xgb_model = XGBClassifier(n_estimators=200, max_depth=4, learning_rate=0.05, use_label_encoder=False, eval_metric='logloss')
xgb_model.fit(X_train, y_train)

xgb_importance = pd.Series(xgb_model.feature_importances_, index=X.columns).sort_values(ascending=False)

plt.figure(figsize=(5, 5))
sns.barplot(x=xgb_importance[:20], y=xgb_importance.index[:20])
plt.title("Top 20 Feature Importances from XGBoost")
plt.xlabel("Importance Score")
plt.tight_layout()
plt.show()



1.0.4 Model Training and Evaluation Using Selected Features

```
[20]: # Top 25 features selected previously
top_xgb_features = ['cloud', 'weather_severity', 'humid_temp_interaction', 'cloud_sunshine_ratio', 'dewpoint',
                     'humid_pressure_interaction', 'pressure_humidity_ratio', 'maxtemp', 'sunshine', 'mintemp',
                     'humid_dewpoint_interaction', 'temp_from_dewpoint', 'day_sin', 'pressure_temp_humidity',
                     'wind_power', 'day_cos', 'temp_range', 'temparature', 'pressure', 'windspeed', 'winddirection',
                     'temp_ratio', 'day', 'max_min_temp_ratio', 'pressure_temp_ratio']

# Prepare inputs
X = train_fe[top_xgb_features]
y = train_fe['rainfall']
X_scaled = MinMaxScaler().fit_transform(X)
```

```

# Class weights
class_weights = compute_class_weight('balanced', classes=np.unique(y), y=y)
weights_dict = dict(zip(np.unique(y), class_weights))

# Initialize models
logreg_model = LogisticRegression(max_iter=1000, class_weight=weights_dict)
lgb_model = LGBMClassifier(n_estimators=100, class_weight='balanced',
                           random_state=42, verbosity=-1)
xgb_model = XGBClassifier(n_estimators=100, use_label_encoder=False,
                           eval_metric='logloss', scale_pos_weight=class_weights[1]/class_weights[0])
stack_model = StackingClassifier(
    estimators=[('lr', Pipeline([('scaler', MinMaxScaler()), ('logreg',
                           LogisticRegression(max_iter=1000, class_weight=weights_dict))]),
                ('lgb', lgb_model), ('xgb', xgb_model)], final_estimator=LogisticRegression())

# Cross-validation
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
results = {'Logistic Regression': [], 'LightGBM': [], 'XGBoost': [], 'CatBoost':
           [], 'Stacking Ensemble': []}

for train_idx, val_idx in skf.split(X, y):
    X_tr, X_val = X.iloc[train_idx], X.iloc[val_idx]
    y_tr, y_val = y.iloc[train_idx], y.iloc[val_idx]

    logreg_model.fit(MinMaxScaler().fit_transform(X_tr), y_tr)
    lgb_model.fit(X_tr, y_tr)
    xgb_model.fit(X_tr, y_tr)
    stack_model.fit(X_tr, y_tr)

    results['Logistic Regression'].append(roc_auc_score(y_val, logreg_model.
                                                       predict_proba(MinMaxScaler().fit_transform(X_val))[:, 1]))
    results['LightGBM'].append(roc_auc_score(y_val, lgb_model.
                                             predict_proba(X_val)[:, 1]))
    results['XGBoost'].append(roc_auc_score(y_val, xgb_model.
                                             predict_proba(X_val)[:, 1]))
    results['Stacking Ensemble'].append(roc_auc_score(y_val, stack_model.
                                                       predict_proba(X_val)[:, 1]))

# Calculate mean AUC for each model
mean_auc_results = {model: np.mean(scores) for model, scores in results.items()}
mean_auc_results

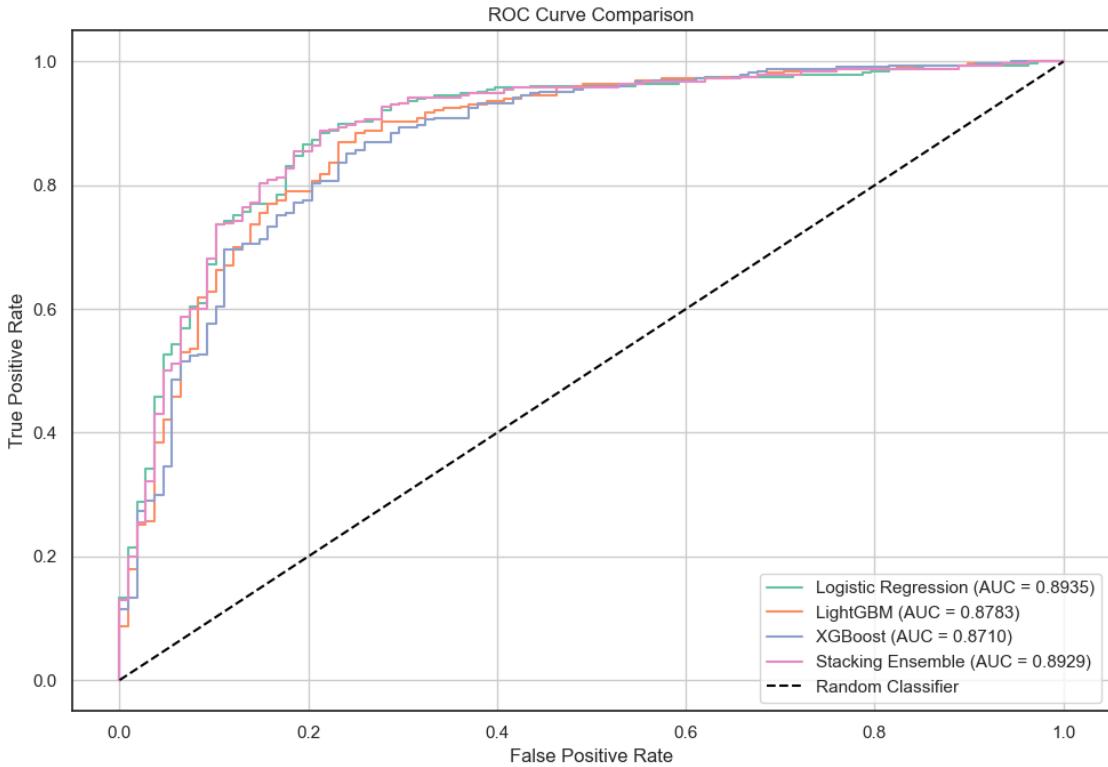
```

[20]: {'Logistic Regression': np.float64(0.8922166105499437),
 'LightGBM': np.float64(0.8716442199775534),

```
'XGBoost': np.float64(0.8661223344556678),  
'CatBoost': np.float64(nan),  
'Stacking Ensemble': np.float64(0.8905667789001122)}
```

1.0.5 ROC Curve Comparison of Classification Models

```
[21]: plt.figure(figsize=(10, 7))  
  
# Plot ROC for each model  
for name, model in [('Logistic Regression', logreg_model),  
                     ('LightGBM', lgb_model),  
                     ('XGBoost', xgb_model),  
                     ('Stacking Ensemble', stack_model)]:  
  
    # Predict probabilities  
    if name == 'Logistic Regression':  
        probas = model.predict_proba(MinMaxScaler().fit_transform(X_val))[:, 1]  
    else:  
        probas = model.predict_proba(X_val)[:, 1]  
  
    fpr, tpr, _ = roc_curve(y_val, probas)  
    roc_auc = auc(fpr, tpr)  
  
    plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.4f})')  
  
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('ROC Curve Comparison')  
plt.legend(loc='lower right')  
plt.grid(True)  
plt.tight_layout()  
plt.show()
```



The ROC curve comparison illustrates the performance of four classification models: Logistic Regression, LightGBM, XGBoost, and a Stacking Ensemble. Among these, **Logistic Regression achieved the highest AUC score of 0.8935**, indicating that the selected features have a strong linear relationship with the target variable. The **Stacking Ensemble closely followed with an AUC of 0.8929**, but did not significantly outperform Logistic Regression, likely due to limited diversity among base models.

Both **LightGBM (AUC = 0.8783)** and **XGBoost (AUC = 0.8710)** also performed well, but slightly lagged behind, suggesting that the additional complexity of these gradient boosting models did not offer substantial gains over the simpler model.

All models significantly outperformed the random classifier baseline, demonstrating strong discriminatory power. Overall, the results highlight the effectiveness of the top 25 selected features and suggest that **Logistic Regression offers a highly effective and interpretable solution for this classification task**.

1.0.6 Final Model Prediction and Submission Preparation

```
[22]: # apply feature engineering to the test set
test_fe = create_features(test_df)

# Select only the top features
X_train_final = train_fe[top_xgb_features]
y_train_final = train_fe['rainfall']
```

```

X_test_final = test_fe[top_xgb_features]

# Normalize features
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train_final)
X_test_scaled = scaler.transform(X_test_final)

# Final Logistic Regression model
final_logreg = LogisticRegression(max_iter=1000, class_weight='balanced')
final_logreg.fit(X_train_scaled, y_train_final)

# Predict probabilities on test set
rainfall_probs = final_logreg.predict_proba(X_test_scaled)[:, 1]

# Prepare submission DataFrame
submission_df = pd.DataFrame({
    'id': test_df['id'],
    'rainfall_probability': rainfall_probs
})

# Save to CSV
submission_path = '/Users/nav/Documents/Jupyter Notebooks /  

    ↴logreg_rainfall_submission.csv'
submission_df.to_csv(submission_path, index=False)

```

1.0.7 Hyperparameter Optimization and Final Submission Using Logistic Regression

[23]:

```

# Scale features
X = train_fe[top_xgb_features]
y = train_fe['rainfall']
X_scaled = MinMaxScaler().fit_transform(X)

def objective(trial):
    C = trial.suggest_loguniform('C', 1e-3, 10)
    penalty = trial.suggest_categorical('penalty', ['l2'])
    solver = trial.suggest_categorical('solver', ['lbfgs', 'liblinear', 'saga'])

    model = LogisticRegression(
        C=C,
        penalty=penalty,
        solver=solver,
        class_weight='balanced',
        max_iter=1000
    )

    return cross_val_score(model, X_scaled, y, cv=5, scoring='roc_auc').mean()

```

```

study = optuna.create_study(direction='maximize', sampler=TPESampler(seed=42))
study.optimize(objective, n_trials=30)

print("Best AUC:", study.best_value)
print("Best Params:", study.best_params)

```

Best AUC: 0.8922109988776654
 Best Params: {'C': 1.4773204665513213, 'penalty': 'l2', 'solver': 'lbfgs'}

```
[24]: # Retrain the final optimized logistic regression model
optimized_logreg = LogisticRegression(
    C=1.4773204665513213,
    penalty='l2',
    solver='lbfgs',
    class_weight='balanced',
    max_iter=1000
)

# Fit on full training set (already filled NaNs and scaled)
X_train_final = train_fe[top_xgb_features]
X_test_final = test_fe[top_xgb_features]
y_train_final = train_fe['rainfall']

scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train_final)
X_test_scaled = scaler.transform(X_test_final)

# Fit and predict
optimized_logreg.fit(X_train_scaled, y_train_final)
rainfall_probs_optimized = optimized_logreg.predict_proba(X_test_scaled)[:, 1]

# Create final submission DataFrame
optimized_submission_df = pd.DataFrame({
    'id': test_df['id'],
    'rainfall_probability': rainfall_probs_optimized})

optimized_submission_path = '/Users/nav/Documents/Jupyter Notebooks /'
optimized_submission_df.to_csv(optimized_submission_path, index=False)
```

[]:

[]:

[]:

[]:

[]: